

WILLIAM WALDEN

wgantt.iv@gmail.com ♦ wgantt.github.io/personal-website ♦ github.com/wgantt

EXPERIENCE

Johns Hopkins University

October 2024 - Present

Research Scientist, [HLTCOE](#) (Primary Appointment)

Assistant Research Scientist, [CLSP](#) (Secondary Appointment)

Baltimore, MD

- Member of the Natural Language Understanding (NLU) team, working on multilingual and multimodal content understanding; evaluation for RAG; and AI for scientific discovery.
- Advise Master's and Ph.D. students in the JHU Center for Language and Speech Processing (CLSP) on research projects in the above areas.
- Founder and lead of the CLSP's AI Safety and Alignment student interest group, facilitating its reading group and advising on related research projects.

Microsoft - Semantic Machines

May 2022 - August 2022

Research Intern

Remote

- Investigated and implemented techniques for calibration and constrained decoding for few-shot semantic parsing with GPT-3 and Codex; improved top- k parsing accuracy by several points absolute on multiple datasets.

Okta

July 2017 - July 2019

Software Engineer

San Francisco, CA

- Led development of Okta's authentication and authorization pipeline, including new frameworks for configurable authorization policies and HTTP callbacks.
- Developed an out-of-the-box self-service registration platform for web apps.

EDUCATION

University of Rochester, Ph.D. Computer Science

June 2024

- Advisor: Aaron Steven White
- Thesis: Document-Level Event Description and Decomposition

University of Rochester, M.S. Computer Science

May 2021

- Advisor: Aaron Steven White

Bowdoin College, B.A. Computer Science (honors), *cum laude*, *Phi Beta Kappa*

May 2017

- Advisors: Clare Bates Congdon and Stephen Majercik
- Thesis: An Investigation of Genetics-Based Machine Learning as Applied to Global Crop Yields

PUBLICATIONS

N.B.: Gantt was my pre-marriage last name. * denotes equal contribution. † denotes equal advising.

Published

- **William Walden***, Kathryn Ricci*, Miriam Wanner, Zhengping Jiang, Chandler May, Rongkun Zhou, Benjamin Van Durme. 2025. [How Grounded is Wikipedia? A Study on Structured Evidential Support and Retrieval](#). *Findings of the Association for Computational Linguistics: ACL 2026*.
- Alexander Martin, Reno Kriz*, **William Walden***, Kate Sanders, Hannah Recknor, Eugene Yang, Francis Ferraro, Benjamin Van Durme. 2025. [WIKIVIDEO: Article Generation from Multiple Videos](#). *Findings of the Association for Computational Linguistics: ACL 2026*.
- **William Walden**, Marc Mason, Orion Weller, Laura Dietz, John Conroy, Neil Molino, Hannah Recknor, Bryan Li, Gabrielle Kaili-May Liu, Yu Hou, James Mayfield, Eugene Yang. 2025. [Auto-ARGUE: LLM-Based Report Generation Evaluation](#). *ACM Special Interest Group on Information Retrieval (SIGIR 2026)*.

- Laura Dietz, Eugene Yang, Bryan Li, Dawn Lawrie, **William Walden**, James Mayfield. 2026. [Insider Knowledge: How Much Can RAG Systems Gain from Evaluation Secrets?](#) *European Conference on Information Retrieval (ECIR 2026)*.
- Laura Dietz, Bryan Li, Gabrielle Kaili Liu, Jia-Huei Ju, Eugene Yang, Dawn Lawrie, **William Walden**, James Mayfield. 2025. [Incorporating Q&A Nuggets into Retrieval-Augmented Generation](#). *European Conference on Information Retrieval (ECIR 2026)*.
- Gabrielle Kaili-May Liu, Bryan Li, Arman Cohan, **William Walden**, Eugene Yang. 2025. [Evaluating Retrieval Augmented Generation Systems on Unanswerable, Uncheatable, Realistic, Multi-hop Queries](#). *European Conference on Information Retrieval (ECIR 2026)*.
- Jiefu Ou*, **William Walden***, Kate Sanders, Zhengping Jiang, Kaiser Sun, Jeffrey Cheng, William Jurayj, Miriam Wanner, Shaobo Liang, Candice Morgan, Seunghoon Han, Weiqi Wang, Chandler May, Hannah Recknor, Daniel Khashabi, Benjamin Van Durme. 2025. [CLAIMCHECK: How Grounded Are LLM Critiques of Scientific Papers?](#). *Findings of the Association for Computational Linguistics: EMNLP 2025*.
- **William Walden**, Pavlo Kuchmiichuk, Alexander Martin, Chihsheng Jin, Angela Cao, Claire Sun, Curisia Allen, Aaron Steven White. 2024. [Cross-Document Event-Keyed Summarization](#). *The First Joint Workshop on Large Language Models and Structure Modeling @ ACL 2025*.
- **William Gantt**. [Document-Level Event Description and Decomposition](#). 2024. *Ph.D. Thesis*.
- **William Gantt**, Alexander Martin, Pavlo Kuchmiichuk, Aaron Steven White. 2024. [Event-Keyed Summarization](#). *Findings of the Association for Computational Linguistics: EMNLP 2024*.
- Siddharth Vashishtha, Alexander Martin, **William Gantt**, Benjamin Van Durme, Aaron Steven White. 2024. [FAMuS: Frames Across Multiple Sources](#). *North American Chapter of the Association for Computational Linguistics (NAACL 2024)*.
- **William Gantt**, Shabnam Behzad, Hannah YoungEun An, Yunmo Chen, Aaron Steven White, Benjamin Van Durme, Mahsa Yarmohammadi. 2024. [MultiMUC: Multilingual Template Filling on MUC-4](#). *European Chapter of the Association for Computational Linguistics (EACL 2024)*.
- **William Gantt**, Reno Kriz*, Yunmo Chen*, Siddharth Vashishtha*, Aaron Steven White. 2023. [On Event Individuation for Document-Level Information Extraction](#). *Findings of the Association for Computational Linguistics: EMNLP 2023*.
- Yunmo Chen*, **William Gantt***, Tongfei Chen*, Aaron Steven White, Benjamin Van Durme. 2023. [A Unified View of Evaluation Metrics for Structured Prediction](#). *Empirical Methods in Natural Language Processing (EMNLP 2023)*.
- Yunmo Chen, **William Gantt**, Weiwei Gu, Tongfei Chen, Aaron Steven White, Benjamin Van Durme. 2023. [Iterative Document-Level Information Extraction via Imitation Learning](#). *European Chapter of the Association for Computational Linguistics (EACL 2022)*. **Outstanding Paper Award**.
- **William Gantt**, Lelia Glass, Aaron Steven White. 2022. [Decomposing and Recomposing Event Structure](#). *Transactions of the Association for Computational Linguistics (TACL)*.
- Benjamin Kane, **William Gantt**, Aaron Steven White. 2021. [Intensional Gaps: Relating doxasticity, bouleticity, veridicality, factivity, and neg-raising](#). *Semantics and Linguistic Theory (SALT 2021)*.
- **William Gantt***, Benjamin Kane*, Aaron Steven White. 2020. [Natural Language Inference with Mixed Effects](#). *The Ninth Joint Conference on Lexical and Computational Semantics (*SEM 2020)*.

In Submission & Preprints

- Miriam Wanner, Hannah Collison, William Jurayj, Benjamin Van Durme, **William Walden**. 2026. Weird Generalization is Weirdly Brittle. *Under Review*.

- Ziyang Huang, Yi Cao, Ali K. Shargh, Jing Luo, Ruidong Mei, Mohd Zaki, Zhan Liu, Wyatt Bunstine, William Jurayj, Somdatta Goswami, Tyrel McQueen, Michael Shields, Jaafar El-Awady, Paulette Clancy, **William Gantt Walden**[†], Nicholas Andrews[†], Benjamin Van Durme, Daniel Khashabi[†]. 2026. RepliCan: Evaluating LLM Agents on Scientific Reproducibility in Computational Materials Science. *Under Review*.
- Bryan Li, **William Walden**, Gabrielle Kaili-May Liu, Yu Hou, Dawn Lawrie, James Mayfield, Eugene Yang, Chris Callison-Burch, Laura Dietz. 2026. DoGMatIQ: Automated Generation of Question-and-Answer Nuggets for Report Evaluation. *Under Review*
- **William Walden** and Miriam Wanner. 2026. [Reasoning Models Will Blatantly Lie About Their Reasoning](#). *Under Review*.
- Yu Hou, Hal Daumé III, Rachel Rudinger, **William Walden**. 2026. On the Role of Citations in Preference Data. *Under Review*.
- Alexander Martin, **William Walden**^{*}, Reno Kriz^{*}, Dengjia Zhang, Chihsheng Jin, Kate Sanders, Eugene Yang, Benjamin Van Durme. 2025. [Seeing through the MiRAGE: Evaluating Multimodal Retrieval Augmented Generation](#). *Under Review*.
- **William Gantt** and Aaron Steven White. 2024. [Small Models Are \(Still\) Effective Cross-Domain Argument Extractors](#). *Preprint*.

SERVICE

Advising & Mentorship

- Miriam Wanner (Johns Hopkins University; Ph.D. expected 2028)
- Ziyang Huang (Johns Hopkins University; M.S. expected 2026)
- Rongkun Zhou (Johns Hopkins University; M.S. 2025)
- Chihsheng Jin (University of Rochester; M.S. 2024)
- Alexander Martin (University of Rochester; B.S. 2024 → Ph.D. student at Johns Hopkins University)
- Weiwei Gu (University of Rochester; M.S. 2022 → Ph.D. student at Arizona State University)

Teaching

- Machine Learning (CSC 246/446): TA, Spring 2021
- Statistical Speech and Language Processing (CSC 248/448): TA, Fall 2020
- Machines and Consciousness (CSC 191/291): TA, Spring 2020

University Service

- JHU North American Computational Linguistics Olympiad Faculty Coordinator: 2025-2026
- URCS Department Graduate Student Representative: 2023-2024
- URCS Department Ph.D. Admissions Committee: 2023-2024

Reviewing

- SIGIR: 2026
- ECIR: 2026
- ACL: 2023-Present
- EACL: 2023, 2024
- EMNLP: 2022-Present
- NAACL: 2021-Present
- ACL Rolling Review: 2021-Present

GRANTS

Scientific Feasibility (SciFy)

DARPA, 2025-2028

- Co-PI for the Johns Hopkins University TA1 team. The SciFy program aims to develop AI tools for automated assessment of the feasibility of novel scientific claims. Award amount: \$3,347,117.

HONORS & AWARDS

AI & National Security Convergence Fellow

Council on Strategic Risks, 2025-2026

- One of 15 fellows selected by the Council on Strategic Risks to participate in a multi-month workshop series on understanding severe risks at the intersection of AI and U.S. National Security.

Outstanding Paper Award

EACL 2024

- For [Iterative Document-Level Information Extraction via Imitation Learning](#)

Outstanding Reviewer

EMNLP 2024

- For high-quality primary reviews of EMNLP 2024 submissions.

Sproull Fellowship

University of Rochester, September 2019

- The University of Rochester's most prestigious graduate fellowship, awarded to fewer than a dozen incoming PhD students on the basis of an outstanding academic record and unusual potential for graduate study.

NSF Research Traineeship

University of Rochester, September 2019

- Full-stipend one-year fellowship awarded to a small set of PhD students in Computer Science and Brain and Cognitive Sciences, focused on computationally-oriented, interdisciplinary research training.

Computer Science Senior-Year Prize

Bowdoin College, May 2017

- Awarded to the student who has achieved the highest distinction in the major program in computer science.

INVITED TALKS

MARE: Automatic Modality-Agnostic Report Evaluation

Lucca, Italy, April 2025

- Eval4Rag Workshop, European Conference on Information Retrieval (ECIR) 2025

Understanding Events in Multimodal Data via Question Answering

Rochester, USA, April 2025

- Fourth Workshop on Processing and Evaluating Event Representations (PEER 2025)

Cross-Document Event-Keyed Summarization

Rochester, USA, April 2025

- Fourth Workshop on Processing and Evaluating Event Representations (PEER 2025)

Event-Keyed Summarization

Ithaca, USA, April 2024

- Third Workshop on Processing and Evaluating Event Representations (PEER 2024)

Structured Representation and Prediction for Document-Level IE

Rochester, USA, April 2023

- Second Workshop on Processing and Evaluating Event Representation (PEER 2023)

Decomposing and Recomposing Event Structure

Ithaca, USA, April 2022

- First Workshop on Processing and Evaluating Event Representations (PEER 2022)

PROJECTS & DATA

[Decompositional Semantics Initiative](#)

- Dataset and toolkit for commonsense semantic annotations and semantic graphs on top of Universal Dependencies on the English Web Treebank.
- Core contributor to version 2.0 of the Decomp Toolkit.
- Lead developer of UDS-EventStructure dataset ([\[1\]](#)).

[MegaIntensionality](#)

- Co-developer of the MegaIntensionality dataset—a large collection of lexically-triggered belief and desire inferences across 725 English clause-embedding verbs ([2]).

IARPA BETTER

- Multilingual information extraction (IE) and retrieval (IR) competition funded by IARPA.
- One of the lead developers on the IE team led by Benjamin Van Durme at Johns Hopkins.
- Led to multiple publications at top NLP venues ([3], [4], [5]).

MultiMUC

- Lead developer of the MultiMUC corpus, a set of translations of the classic MUC-4 template filling dataset into Arabic, Farsi, Mandarin, Korean, and Russian.
- The only publicly available multilingual template filling corpus ([6]).

MUCSUM

- Lead developer of the MUCSUM dataset, a collection of summaries of all events in the MUC-4 corpus ([7]).

SEAMuS

- Lead developer of the SEAMuS dataset, a collection of single- and cross-document summaries based on the FAMuS dataset for cross-document argument extraction ([8]).

DARPA SciFy

- Program focused on assessing the feasibility of novel claims in different scientific domains. Data and subject matter expert (SME) lead for the JHU team. Resulted in ([9]).

SKILLS

Programming Languages	Python (expert) Java, R, SQL, C, C++, Bash (some experience)
Tools & Libraries	Claude Code, NumPy, Pandas, PyTorch, HuggingFace, Amazon Mechanical Turk, Git, Vim